

black-friday-sales-basic-eda

June 5, 2024

1 Analyzing the Raw data given wrt Black Friday Sales

2 Installing necessary libraries

```
[1]: pip install pandas numpy matplotlib seaborn
```

```
Requirement already satisfied: pandas in  
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages (2.2.2)  
Requirement already satisfied: numpy in  
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages  
(1.26.4)  
Requirement already satisfied: matplotlib in  
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages (3.9.0)  
Requirement already satisfied: seaborn in  
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages  
(0.13.2)  
Requirement already satisfied: python-dateutil>=2.8.2 in  
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages (from  
pandas) (2.9.0.post0)  
Requirement already satisfied: pytz>=2020.1 in  
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages (from  
pandas) (2024.1)  
Requirement already satisfied: tzdata>=2022.7 in  
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages (from  
pandas) (2024.1)  
Requirement already satisfied: contourpy>=1.0.1 in  
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages (from  
matplotlib) (1.2.1)  
Requirement already satisfied: cycler>=0.10 in  
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages (from  
matplotlib) (0.12.1)  
Requirement already satisfied: fonttools>=4.22.0 in  
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages (from  
matplotlib) (4.53.0)  
Requirement already satisfied: kiwisolver>=1.3.1 in  
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages (from  
matplotlib) (1.4.5)
```

Requirement already satisfied: packaging>=20.0 in
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages (from
matplotlib) (24.0)
Requirement already satisfied: pillow>=8 in
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages (from
matplotlib) (10.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages (from
matplotlib) (3.1.2)
Requirement already satisfied: six>=1.5 in
c:\users\shubh\appdata\local\programs\python\python312\lib\site-packages (from
python-dateutil>=2.8.2->pandas) (1.16.0)
Note: you may need to restart the kernel to use updated packages.

```
[2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

3 Reading the csv file

```
[3]: df = pd.read_csv(r'C:\Users\Shubh\Desktop\PROJECTS\PYTHON\EDA_
↳_BlackFridaySales\Data\BlackFriday.csv')
```

4 Total number of rows and column

To find total number of rows and column, here we have stored csv file in “df” variable. If we execute this df we will get total rows and total columns

Here total rows : 537577 rows total columns : 12 columns

```
[4]: df
```

```
[4]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	\
0	1000001	P00069042	F	0-17	10	A	
1	1000001	P00248942	F	0-17	10	A	
2	1000001	P00087842	F	0-17	10	A	
3	1000001	P00085442	F	0-17	10	A	
4	1000002	P00285442	M	55+	16	C	
...	
537572	1004737	P00193542	M	36-45	16	C	
537573	1004737	P00111142	M	36-45	16	C	
537574	1004737	P00345942	M	36-45	16	C	
537575	1004737	P00285842	M	36-45	16	C	
537576	1004737	P00118242	M	36-45	16	C	

	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	\
0	2	0	3	
1	2	0	1	
2	2	0	12	
3	2	0	12	
4	4+	0	8	
...	
537572	1	0	1	
537573	1	0	1	
537574	1	0	8	
537575	1	0	5	
537576	1	0	5	

	Product_Category_2	Product_Category_3	Purchase
0	NaN	NaN	8370
1	6.0	14.0	15200
2	NaN	NaN	1422
3	14.0	NaN	1057
4	NaN	NaN	7969
...
537572	2.0	NaN	11664
537573	15.0	16.0	19196
537574	15.0	NaN	8043
537575	NaN	NaN	7172
537576	8.0	NaN	6875

[537577 rows x 12 columns]

5 Analyzing Dataframe (initial level)

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 537577 entries, 0 to 537576
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   User_ID               537577 non-null  int64
1   Product_ID            537577 non-null  object
2   Gender                537577 non-null  object
3   Age                  537577 non-null  object
4   Occupation            537577 non-null  int64
5   City_Category        537577 non-null  object
6   Stay_In_Current_City_Years  537577 non-null  object
7   Marital_Status        537577 non-null  int64
8   Product_Category_1    537577 non-null  int64
```

```

9   Product_Category_2      370591 non-null  float64
10  Product_Category_3      164278 non-null  float64
11  Purchase                537577 non-null  int64

```

dtypes: float64(2), int64(5), object(5)

memory usage: 49.2+ MB

Observation 1: Now here we can notice age datatype is object which should be int ;

In Product_Category_2 and Product_Category_3 there seems to be some null values
rest all seems ok.

[6]: df

```

[6]:      User_ID Product_ID Gender   Age Occupation City_Category \
0      1000001  P00069042      F  0-17          10           A
1      1000001  P00248942      F  0-17          10           A
2      1000001  P00087842      F  0-17          10           A
3      1000001  P00085442      F  0-17          10           A
4      1000002  P00285442      M   55+          16           C
...      ...      ...      ...      ...      ...      ...
537572  1004737  P00193542      M  36-45          16           C
537573  1004737  P00111142      M  36-45          16           C
537574  1004737  P00345942      M  36-45          16           C
537575  1004737  P00285842      M  36-45          16           C
537576  1004737  P00118242      M  36-45          16           C

      Stay_In_Current_City_Years  Marital_Status  Product_Category_1 \
0                                2                0                3
1                                2                0                1
2                                2                0               12
3                                2                0               12
4                                4+                0                8
...      ...      ...      ...      ...
537572                        1                0                1
537573                        1                0                1
537574                        1                0                8
537575                        1                0                5
537576                        1                0                5

      Product_Category_2  Product_Category_3  Purchase
0                    NaN                  NaN       8370
1                    6.0                  14.0      15200
2                    NaN                  NaN       1422
3                   14.0                  NaN       1057
4                    NaN                  NaN       7969
...      ...      ...      ...
537572                    2.0                  NaN      11664
537573                   15.0                  16.0      19196
537574                   15.0                  NaN       8043

```

537575	NaN	NaN	7172
537576	8.0	NaN	6875

[537577 rows x 12 columns]

Identifying columns with missing values and decide how to handle them, such as imputing missing

6 Handling Missing Value

The `df.isnull().sum()` method in pandas is used to count the number of missing values (NaN) in each column of a DataFrame.

Here's a breakdown of how it works:

- 1) `df.isnull()` returns a DataFrame of the same shape as the original DataFrame (`df`), where each cell contains a boolean value indicating whether the corresponding cell in the original DataFrame is missing (True) or not (False).
- 2) `sum()` is then applied to this DataFrame, which sums up the boolean values along each column. Since in Python, True is treated as 1 and False as 0 when summed, this effectively counts the number of missing values in each column.

When we run `df.isnull().sum()`, we will get a Series object where each index corresponds to a column name in `df`, and each value indicates the number of missing values in that column.

```
[7]: df.isnull().sum()
```

```
[7]: User_ID          0
     Product_ID      0
     Gender          0
     Age             0
     Occupation      0
     City_Category   0
     Stay_In_Current_City_Years  0
     Marital_Status  0
     Product_Category_1  0
     Product_Category_2 166986
     Product_Category_3 373299
     Purchase         0
     dtype: int64
```

Observation2 : Product_Category_2 and Product_Category_3 have 166986 and 373299 missing values

The `dropna()` method is commonly used in data preprocessing to remove rows or columns with missing data before analysis. The `df.dropna()` method in pandas is used to remove rows or columns from a DataFrame that contain any missing values (NaN).

```
#Remove rows with any missing values df_cleaned = df.dropna()
```

```
#Remove columns with any missing values df_cleaned = df.dropna(axis=1)
```

```
#Remove rows with missing values in specific columns df_cleaned = df.dropna(subset=['column1',
'column2'])
```

```
#Remove rows with at least 3 non-missing values df_cleaned = df.dropna(thresh=3)
```

7 Removing Missing Values in Rows and Column

```
[8]: df.dropna()
```

```
[8]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	\
1	1000001	P00248942	F	0-17	10	A	
6	1000004	P00184942	M	46-50	7	B	
13	1000005	P00145042	M	26-35	20	A	
14	1000006	P00231342	F	51-55	9	A	
16	1000006	P0096642	F	51-55	9	A	
...	
537549	1004734	P00345842	M	51-55	1	B	
537551	1004735	P00313442	M	46-50	3	C	
537562	1004736	P00146742	M	18-25	20	A	
537571	1004737	P00221442	M	36-45	16	C	
537573	1004737	P00111142	M	36-45	16	C	

	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	\
1		2	0	1
6		2	1	1
13		1	1	1
14		1	0	5
16		1	0	2
...
537549		1	1	2
537551		3	0	5
537562		1	1	1
537571		1	0	1
537573		1	0	1

	Product_Category_2	Product_Category_3	Purchase
1	6.0	14.0	15200
6	8.0	17.0	19215
13	2.0	5.0	15665
14	8.0	14.0	5378
16	3.0	4.0	13055
...
537549	8.0	14.0	13082
537551	6.0	8.0	6863
537562	13.0	14.0	11508
537571	2.0	5.0	11852
537573	15.0	16.0	19196

[164278 rows x 12 columns]

Writing below

```
del df['Product_Category_2'] del df['Product_Category_3']
```

or `df.dropna(axis=1)` will do one and the same thing as here we need to delete the entire column.

```
[9]: df.dropna(axis=1)
```

```
[9]:
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	\
0	1000001	P00069042	F	0-17	10	A	
1	1000001	P00248942	F	0-17	10	A	
2	1000001	P00087842	F	0-17	10	A	
3	1000001	P00085442	F	0-17	10	A	
4	1000002	P00285442	M	55+	16	C	
...	
537572	1004737	P00193542	M	36-45	16	C	
537573	1004737	P00111142	M	36-45	16	C	
537574	1004737	P00345942	M	36-45	16	C	
537575	1004737	P00285842	M	36-45	16	C	
537576	1004737	P00118242	M	36-45	16	C	

	Stay_In_Current_City_Years	Marital_Status	Product_Category_1	\
0	2	0	3	
1	2	0	1	
2	2	0	12	
3	2	0	12	
4	4+	0	8	
...	
537572	1	0	1	
537573	1	0	1	
537574	1	0	8	
537575	1	0	5	
537576	1	0	5	

	Purchase
0	8370
1	15200
2	1422
3	1057
4	7969
...	...
537572	11664
537573	19196
537574	8043
537575	7172

537576 6875

[537577 rows x 10 columns]

[10]: df

```
[10]:      User_ID Product_ID Gender   Age Occupation City_Category \
0      1000001  P00069042      F  0-17          10           A
1      1000001  P00248942      F  0-17          10           A
2      1000001  P00087842      F  0-17          10           A
3      1000001  P00085442      F  0-17          10           A
4      1000002  P00285442      M   55+          16           C
...      ...      ...      ...      ...      ...      ...
537572  1004737  P00193542      M  36-45          16           C
537573  1004737  P00111142      M  36-45          16           C
537574  1004737  P00345942      M  36-45          16           C
537575  1004737  P00285842      M  36-45          16           C
537576  1004737  P00118242      M  36-45          16           C

      Stay_In_Current_City_Years  Marital_Status  Product_Category_1 \
0                                2                0                3
1                                2                0                1
2                                2                0               12
3                                2                0               12
4                                4+                0                8
...      ...      ...      ...      ...
537572                                1                0                1
537573                                1                0                1
537574                                1                0                8
537575                                1                0                5
537576                                1                0                5

      Product_Category_2  Product_Category_3  Purchase
0                    NaN                    NaN      8370
1                    6.0                   14.0     15200
2                    NaN                    NaN      1422
3                   14.0                    NaN      1057
4                    NaN                    NaN      7969
...      ...      ...      ...
537572                    2.0                    NaN     11664
537573                    15.0                   16.0     19196
537574                    15.0                    NaN      8043
537575                    NaN                    NaN      7172
537576                    8.0                    NaN      6875
```

[537577 rows x 12 columns]

8 Step 2 : Now Lets Analyze the columns of the data

```
[11]: df.head()
```

```
[11]:   User_ID Product_ID Gender   Age Occupation City_Category \
0  1000001  P00069042      F  0-17          10           A
1  1000001  P00248942      F  0-17          10           A
2  1000001  P00087842      F  0-17          10           A
3  1000001  P00085442      F  0-17          10           A
4  1000002  P00285442      M  55+          16           C

   Stay_In_Current_City_Years  Marital_Status  Product_Category_1 \
0                             2                0                  3
1                             2                0                  1
2                             2                0                 12
3                             2                0                 12
4                             4+                0                  8

   Product_Category_2  Product_Category_3  Purchase
0                  NaN                  NaN       8370
1                   6.0                  14.0      15200
2                  NaN                  NaN       1422
3                  14.0                  NaN       1057
4                  NaN                  NaN       7969
```

```
[12]: df['User_ID'].nunique()
```

```
[12]: 5891
```

```
[13]: df['Product_ID'].nunique()
```

```
[13]: 3623
```

```
[14]: df['Gender'].unique()
```

```
[14]: array(['F', 'M'], dtype=object)
```

```
[15]: df['Age'].unique()
```

```
[15]: array(['0-17', '55+', '26-35', '46-50', '51-55', '36-45', '18-25'],
      dtype=object)
```

```
[16]: df['Occupation'].unique()
```

```
[16]: array([10, 16, 15,  7, 20,  9,  1, 12, 17,  0,  3,  4, 11,  8, 19,  2, 18,
        5, 14, 13,  6], dtype=int64)
```

```
[17]: df['City_Category'].unique()

[17]: array(['A', 'C', 'B'], dtype=object)

[18]: df['Stay_In_Current_City_Years'].unique()

[18]: array(['2', '4+', '3', '1', '0'], dtype=object)

[19]: df['Marital_Status'].unique()

[19]: array([0, 1], dtype=int64)

[20]: df['Product_Category_1'].unique()

[20]: array([ 3,  1, 12,  8,  5,  4,  2,  6, 14, 11, 13, 15,  7, 16, 18, 10, 17,
           9], dtype=int64)

[21]: df['Purchase'].sum()/len(df['Purchase'])

[21]: 9333.859852635065

[22]: for column in df.columns:
      print(column, ":", df[column].nunique())
```

```
User_ID : 5891
Product_ID : 3623
Gender : 2
Age : 7
Occupation : 21
City_Category : 3
Stay_In_Current_City_Years : 5
Marital_Status : 2
Product_Category_1 : 18
Product_Category_2 : 17
Product_Category_3 : 15
Purchase : 17959
```

9 Step3: Analyzing Gender

```
[23]: # Create a DataFrame to hold unique counts of each column
unique_counts_df = pd.DataFrame({
    'Column': df.columns,
    'Unique Counts': [df[column].nunique() for column in df.columns]
})

# Display the unique counts DataFrame in tabular form
display(unique_counts_df)
```


	Column	Unique Counts
0	User_ID	5891
1	Product_ID	3623
2	Gender	2
3	Age	7
4	Occupation	21
5	City_Category	3
6	Stay_In_Current_City_Years	5
7	Marital_Status	2
8	Product_Category_1	18
9	Product_Category_2	17
10	Product_Category_3	15
11	Purchase	17959

Creating a pie chart to visualize the gender distribution in the DataFrame df.

```
[24]: data = pd.DataFrame({'Ratio' : [len(df[df['Gender'] == 'M']),
    ↪len(df[df['Gender'] == 'F'])]},
    index = ['Male', 'Female'])

data.plot.pie(y = 'Ratio', figsize = (6,6), autopct = "%.1f")
```


```
[24]: <Axes: ylabel='Ratio'>
```



```
[25]: df.groupby('Gender').size().plot(kind = 'pie',  
                                         autopct = "%.1f",  
                                         title = 'Gender Ratio',  
                                         figsize = (6,6))
```


```
[25]: <Axes: title={'center': 'Gender Ratio'}>
```

Gender Ratio



```
[26]: df.groupby('Gender').size().plot(kind = 'bar',  
                                         figsize = (6,6))
```

```
[26]: <Axes: xlabel='Gender'>
```




```
[27]: df.groupby('Gender').size()
```

```
[27]: Gender
F      132197
M      405380
dtype: int64
```

10 Analyzing Age and Marital Status

```
[28]: df.groupby('Age').size().plot(kind = 'bar', figsize = (12, 6), title = 'Purchase Distribution by Age')
```


```
[28]: <Axes: title={'center': 'Purchase Distribution by Age'}, xlabel='Age'>
```



```
[29]: lst = []  
      for i in df['Age'].unique():  
          lst.append([i, df[df['Age'] == i]['Product_ID'].nunique()])  
  
      data = pd.DataFrame(lst , columns = ['Age', 'Products'])
```


```
[30]: data.plot.bar(x = 'Age', figsize = (8,6))
```

```
[30]: <Axes: xlabel='Age'>
```



```
[31]: df.groupby('Age').sum()['Purchase'].plot(kind = 'bar', figsize = (12, 6), title='Amount Spend by Age')
```

```
[31]: <Axes: title={'center': 'Amount Spend by Age'}, xlabel='Age'>
```

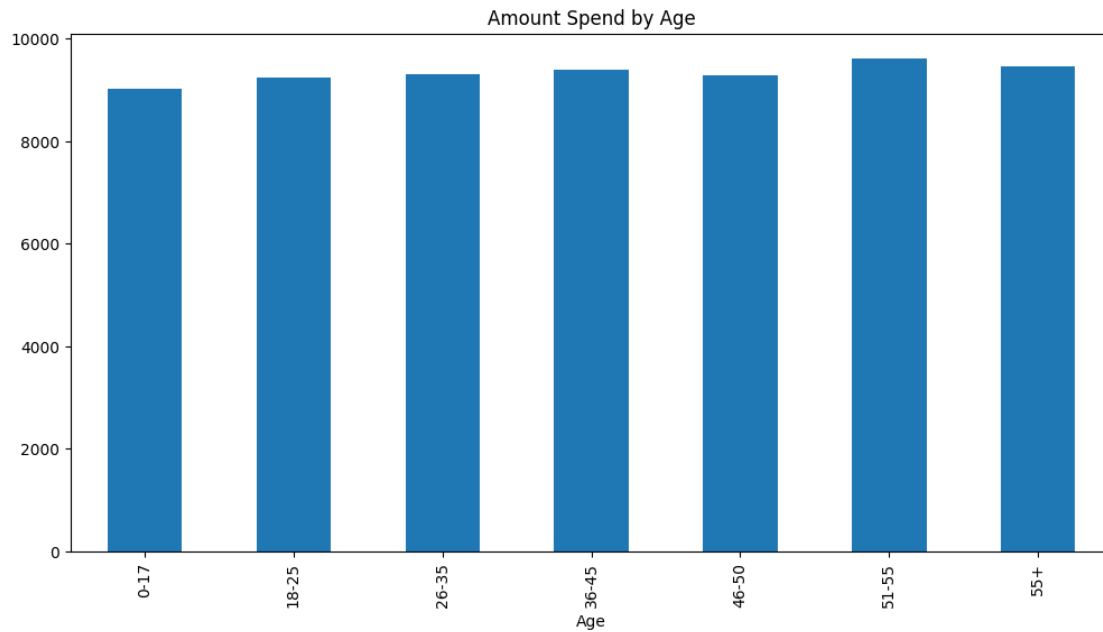
```
[37]: ##df.groupby('Age').mean()['Purchase'].plot(kind = 'bar', figsize = (12, 6),
      ↪title = 'Amount Spend by Age')

# Convert 'Purchase' column to numeric, coercing errors to NaN
df['Purchase'] = pd.to_numeric(df['Purchase'], errors='coerce')

# Drop rows with NaN values in 'Purchase' column
df.dropna(subset=['Purchase'], inplace=True)

# Now, try plotting again
df.groupby('Age')['Purchase'].mean().plot(kind='bar', figsize=(12, 6),
      ↪title='Amount Spend by Age')
```

```
[37]: <Axes: title={'center': 'Amount Spend by Age'}, xlabel='Age'>
```




```
[39]: #df.groupby('Age').mean()['Purchase'].plot(kind = 'pie', autopct = '%0.1f')
```

```
# Group by 'Age' and calculate the mean purchase amount
mean_purchase_by_age = df.groupby('Age')['Purchase'].mean()


# Plot the pie chart
mean_purchase_by_age.plot(kind='pie', autopct='%0.1f%%')
plt.ylabel('') # Remove the y-label
plt.title('Average Purchase Amount by Age Group')
plt.show()
```

Average Purchase Amount by Age Group




```
[40]: df.groupby('Marital_Status').size().plot(kind = 'pie', autopct = '%0.1f')
```

```
[40]: <Axes: >
```



```
[41]: df.groupby('Gender').size().plot(kind = 'pie', autopct = '%0.1f')
```

```
[41]: <Axes: >
```




11 MultiColumn Analysis

```
[42]: import seaborn as sns
```


```
[43]: df.groupby('Age').size().plot(kind = 'bar', figsize = (12, 6), title = 'Purchase Distribution by Age')
```

```
[43]: <Axes: title={'center': 'Purchase Distribution by Age'}, xlabel='Age'>
```




```
[44]: sns.set(rc = {'figure.figsize' : (12,6)})
      sns.countplot(x = "Age", hue = 'Gender', data = df)
```

```
[44]: <Axes: xlabel='Age', ylabel='count'>
```



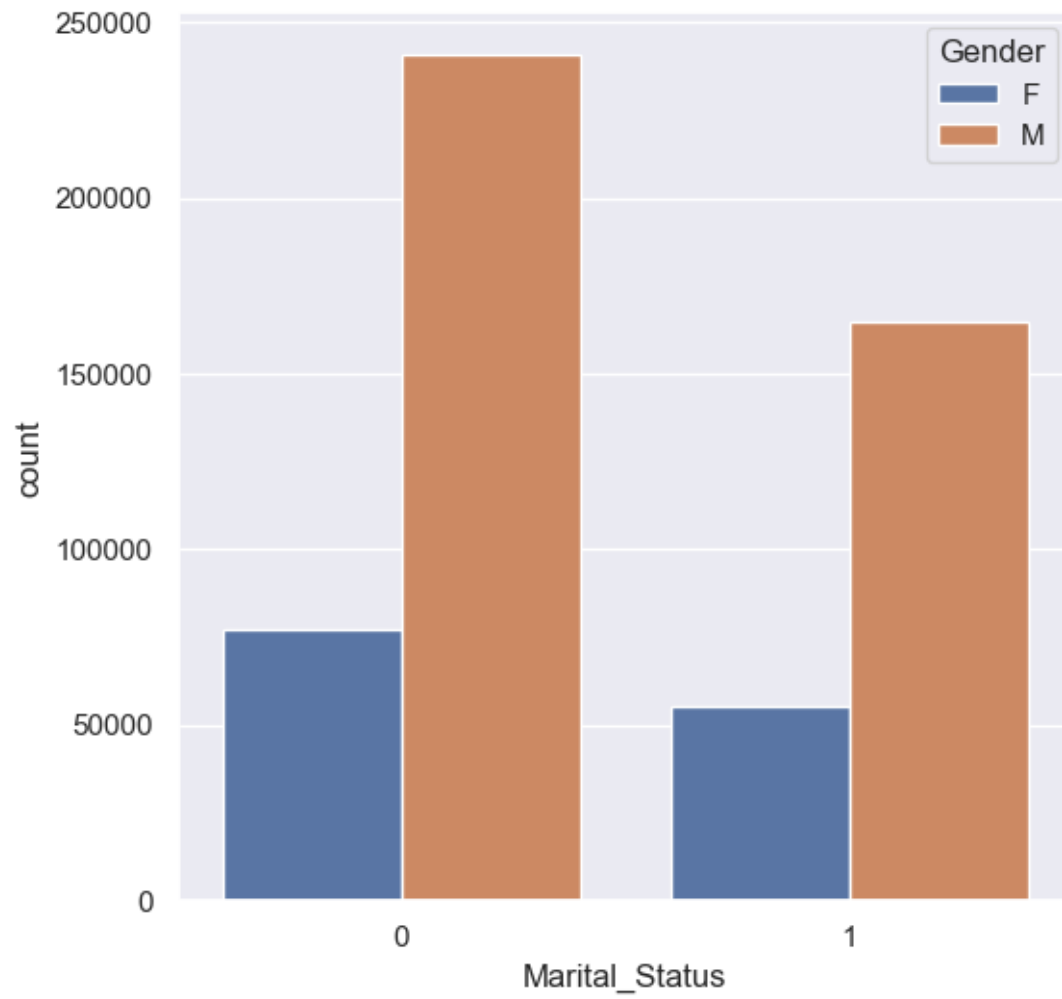
```
[45]: sns.set(rc = {'figure.figsize' : (12,6)})
      sns.countplot(x = "Gender", hue = 'Age', data = df)
```

[45]: <Axes: xlabel='Gender', ylabel='count'>




```
[46]: sns.set(rc = {'figure.figsize' : (6,6)})  
sns.countplot(x = "Marital_Status", hue = 'Gender', data = df)
```

[46]: <Axes: xlabel='Marital_Status', ylabel='count'>




```
[47]: sns.countplot(x = df['City_Category'])
```

```
[47]: <Axes: xlabel='City_Category', ylabel='count'>
```


```
[48]: sns.countplot(x = 'City_Category', hue = 'Age', data = df)
```

```
[48]: <Axes: xlabel='City_Category', ylabel='count'>
```




```
[49]: sns.countplot(x = 'Age', hue = 'City_Category', data = df)
```

```
[49]: <Axes: xlabel='Age', ylabel='count'>
```




```
[50]: sns.countplot(x = 'Marital_Status', hue = 'City_Category', data = df)
```

```
[50]: <Axes: xlabel='Marital_Status', ylabel='count'>
```



```
[51]: df.groupby('City_Category').size().plot(kind = 'pie', autopct = '%0.1f')
```

```
[51]: <Axes: >
```




```
[52]: sns.countplot(x = 'City_Category', hue = 'Gender', data = df)
```

```
[52]: <Axes: xlabel='City_Category', ylabel='count'>
```



```
[53]: df.groupby('City_Category').sum()['Purchase'].plot(kind = 'pie', autopct = "%0.1f")
```


```
[53]: <Axes: ylabel='Purchase'>
```



```
[55]: # Group by 'City_Category' and calculate the mean purchase amount
mean_purchase_by_city = df.groupby('City_Category')['Purchase'].mean()

# Plot the pie chart
mean_purchase_by_city.plot(kind='pie', autopct='%0.1f%%')
plt.ylabel('') # Remove the y-label
plt.title('Average Purchase Amount by City Category')
plt.show()
```

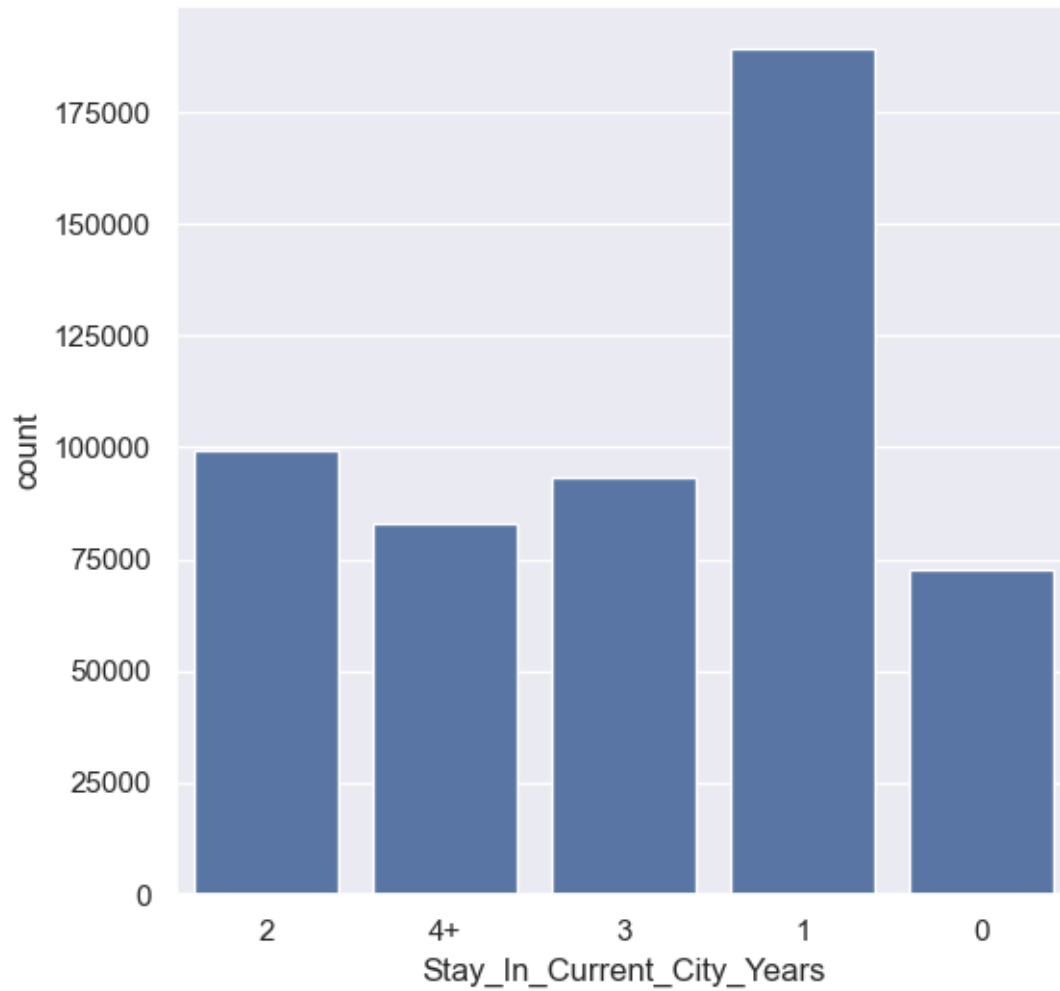
Average Purchase Amount by City Category



12 Occupation and Products Analysis


```
[56]: sns.countplot(x = df['Stay_In_Current_City_Years'])
```

```
[56]: <Axes: xlabel='Stay_In_Current_City_Years', ylabel='count'>
```


```
[57]: df.groupby('Stay_In_Current_City_Years').size().plot(kind = 'pie', autopct = "%.  
      ↪1f")
```

```
[57]: <Axes: >
```



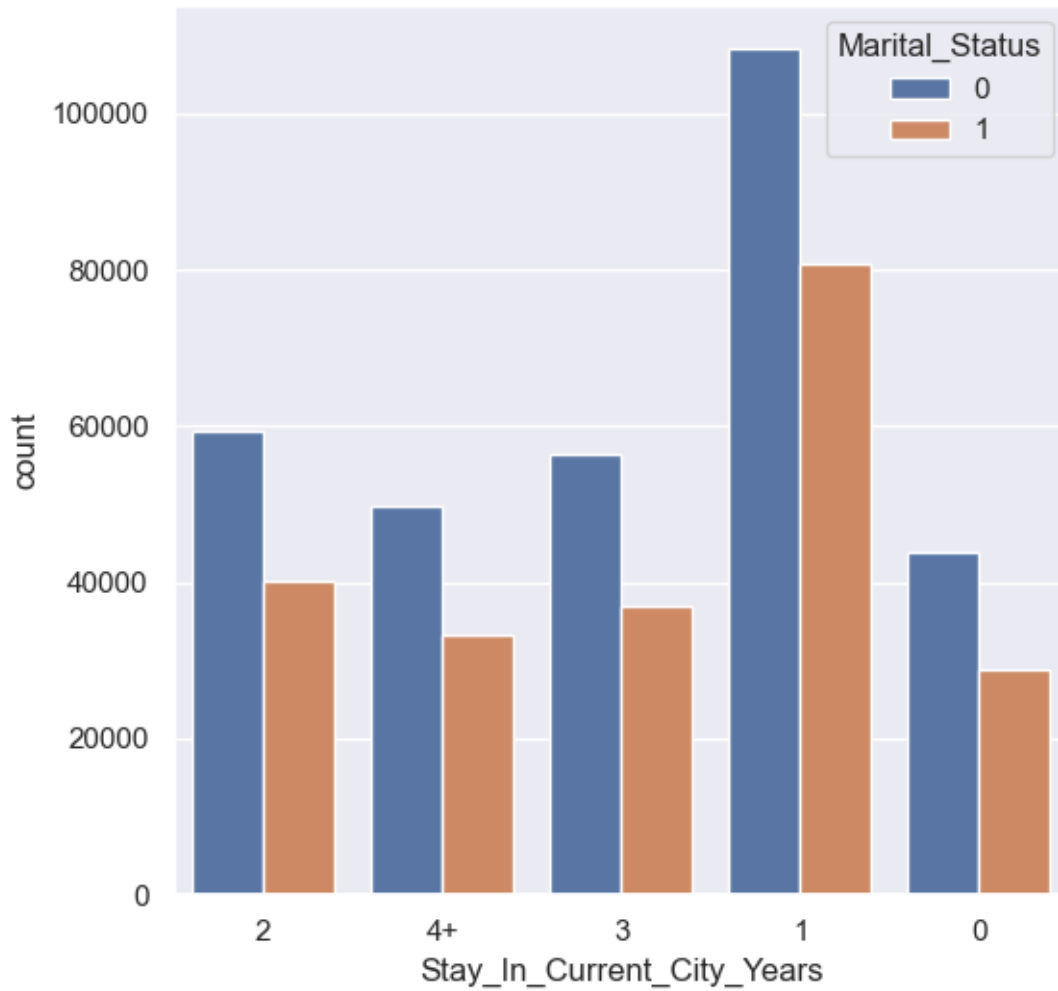
```
[58]: sns.countplot(x = 'Stay_In_Current_City_Years', hue = 'Gender', data = df)
```

```
[58]: <Axes: xlabel='Stay_In_Current_City_Years', ylabel='count'>
```




```
[59]: sns.countplot(x = 'Stay_In_Current_City_Years', hue = 'Marital_Status', data = df)
```

```
[59]: <Axes: xlabel='Stay_In_Current_City_Years', ylabel='count'>
```



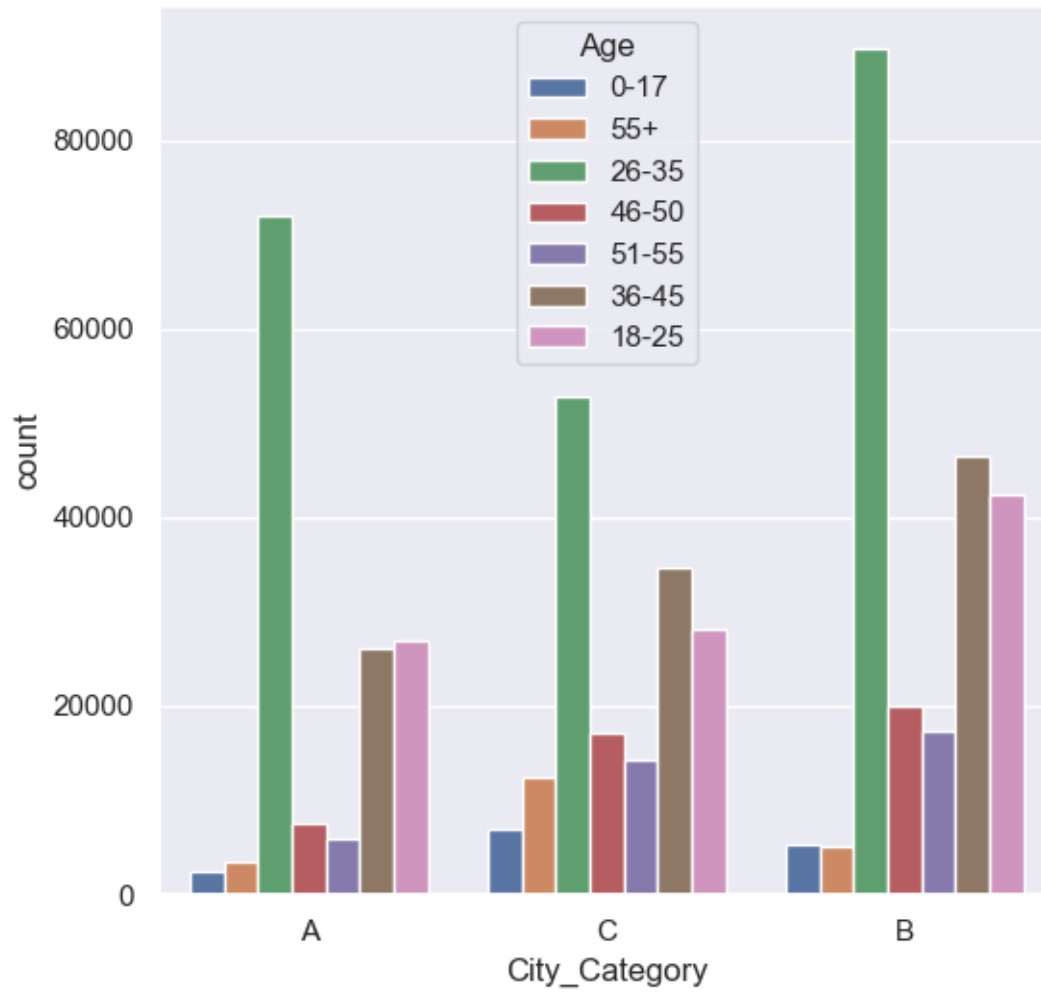
```
[60]: sns.countplot(x = 'Stay_In_Current_City_Years', hue = 'City_Category', data = df)
```

```
[60]: <Axes: xlabel='Stay_In_Current_City_Years', ylabel='count'>
```




```
[61]: sns.countplot(x = 'City_Category', hue = 'Age', data = df)
```

```
[61]: <Axes: xlabel='City_Category', ylabel='count'>
```




```
[63]: df.groupby('Stay_In_Current_City_Years').sum()['Purchase'].plot(kind = 'bar')
```

```
[63]: <Axes: xlabel='Stay_In_Current_City_Years'>
```




```
[67]: sns.countplot(x = df['Occupation'])
```

```
[67]: <Axes: xlabel='Occupation', ylabel='count'>
```




```
[68]: df.groupby('Occupation').size().sort_values().plot(kind = 'bar')
```

```
[68]: <Axes: xlabel='Occupation'>
```

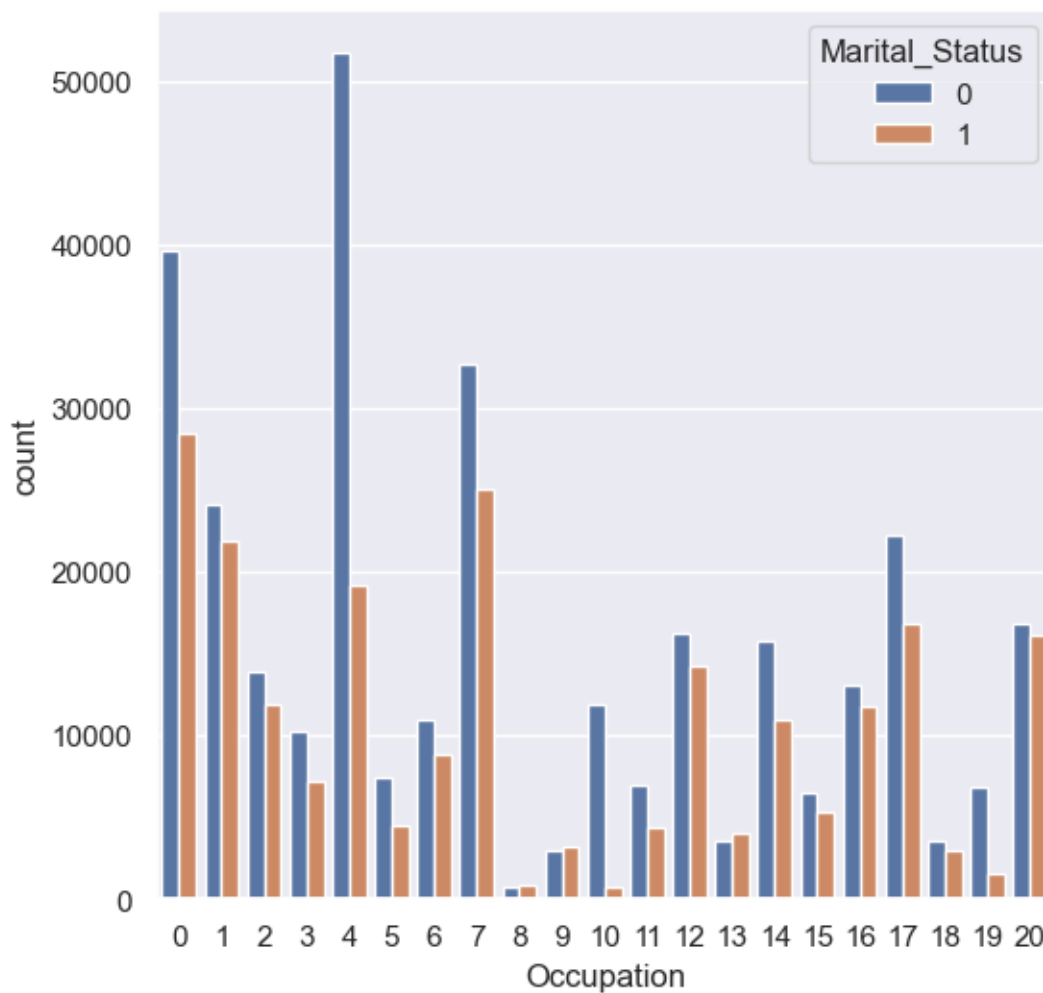
```
[69]: df.groupby('Occupation').sum()['Purchase'].sort_values().plot(kind = 'bar')
```

```
[69]: <Axes: xlabel='Occupation'>
```




```
[74]: sns.countplot(x = 'Occupation', hue = 'Marital_Status', data = df)
```

```
[74]: <Axes: xlabel='Occupation', ylabel='count'>
```




```
[75]: sns.countplot(x = 'Occupation', hue = 'Gender', data = df)
```

```
[75]: <Axes: xlabel='Occupation', ylabel='count'>
```




```
[76]: df.groupby('Occupation').nunique()['Product_ID'].plot(kind = 'bar')
```

```
[76]: <Axes: xlabel='Occupation'>
```




```
[77]: df.groupby('Occupation').nunique()['Product_ID'].sort_values().plot(kind='bar')
```

```
[77]: <Axes: xlabel='Occupation'>
```




```
[78]: df.groupby('Product_Category_1').size().plot(kind = 'bar')
```

```
[78]: <Axes: xlabel='Product_Category_1'>
```




```
[79]: df.groupby('Product_Category_1').size().sort_values().plot(kind = 'bar')
```

```
[79]: <Axes: xlabel='Product_Category_1'>
```




```
[80]: df.groupby('Product_ID').sum()['Purchase'].nlargest(10).sort_values().plot(kind='bar')
```

```
[80]: <Axes: xlabel='Product_ID'>
```


```
[81]: df.groupby('Product_ID').size().nlargest(10).sort_values().plot(kind = 'bar')
```

```
[81]: <Axes: xlabel='Product_ID'>
```




```
[83]: sns.countplot(x = 'Product_Category_1', hue = 'Gender', data = df)
```

```
[83]: <Axes: xlabel='Product_Category_1', ylabel='count'>
```



```
[84]: sns.countplot(x = 'Product_Category_1', hue = 'Marital_Status', data = df)
```

```
[84]: <Axes: xlabel='Product_Category_1', ylabel='count'>
```




13 Combining Age & Marital Status

```
[85]: l = []
      for i in range(len(df)):
          l.append(df['Gender'][i] + "_" + str(df['Marital_Status'][i]))

      df['MaritalGender'] = l
```


```
[86]: sns.countplot(x = df['MaritalGender'])
```

```
[86]: <Axes: xlabel='MaritalGender', ylabel='count'>
```




```
[87]: sns.countplot(x = df['Age'], hue = df['MaritalGender'])
```

```
[87]: <Axes: xlabel='Age', ylabel='count'>
```




```
[88]: sns.countplot(x = df['Product_Category_1'], hue = df['MaritalGender'])
```

```
[88]: <Axes: xlabel='Product_Category_1', ylabel='count'>
```




```
[89]: sns.countplot(x = df['Stay_In_Current_City_Years'], hue = df['MaritalGender'])
```

```
[89]: <Axes: xlabel='Stay_In_Current_City_Years', ylabel='count'>
```



```
[90]: sns.countplot(x = df['City_Category'], hue = df['MaritalGender'])
```

```
[90]: <Axes: xlabel='City_Category', ylabel='count'>
```

[]: